OpenScientists

An open platform for AI-assisted scientific discovery

PROJECT BRIEF · APRIL 2026 · NON-CONFIDENTIAL

> *Ship a working technique directly to the people who use it. Skip the paper-shaped detour.*

The problem

The path from a working technique to the users who benefit from it is too long.

- A researcher who discovers something that works cannot simply share it. They must package it as a paper, hide the simplest piece in a larger narrative, cherry-pick results, and wait on committees.
- With a paper in hand, they approach investors — another round of story-crafting, another set of gatekeepers.
- Customers then demand still more papers before adopting the technique in practice.

Each step was useful once. In the current era of capable AI systems, most add latency rather than value.

Our answer

A point-to-point, public-interest platform where the unit of exchange is not a paper but a **skill** — a reusable, composable agent workflow that performs a concrete scientific task.

Contributed directly. Evaluated openly. Used immediately. Attributed transparently.

Sits inside a larger open agent ecosystem

a general-purpose agent framework, an auto-research framework on top, a scientific skill benchmark, plus harness, memory, agentic RL, and reliable-agent modules — each open source, each with a named maintainer.

What we will do

- Build and operate the open skill platform: share, discover, evaluate, exchange, compose.
- Run empirical studies of AI-mediated contribution in peer review and on public software infrastructure (agent-proposed and agent-reviewed PRs).
- Release open-source modules for the supporting stack.
- Publish a methodology article with senior academic collaborators.
- Maintain and iterate as the primitives of AI-assisted work evolve.

Milestones

NEAR-TERM

Public repository and landing page. First sponsor and partner conversations closed. Non-profit entity formed.

MID-TERM

Platform v0. Initial skill catalog. Seeded benchmark suite. First pilot partners.

LONG-TERM

Self-sustaining open ecosystem. Independent module maintainers. Persistent evaluation. Active community.

What we are asking for

Every category below unlocks work that is otherwise blocked. All funds are received by the non-profit, never by individuals.

Philanthropic funding	Engineering, fellowships, infrastructure, long-term maintenance.
Framework contributions	Auto-research frameworks and agent-startup partnerships that coproduce pipelines and share infrastructure.
Skills & expertise	Initial focus: AI-for-science verticals.
Review & governance	Senior academics, publishers, and scientific societies.
Model-provider credits	Donated, or at cost. Cost-price credits also let us publish transparent benchmarks on the true operating cost of frontier-class agent systems.
Compute & GPUs	For training-heavy modules: agentic RL, reliable-agent work.
Benchmark tasks	Hard, faithful problems encoded by domain researchers.

How contributors are credited

Every contribution is attributed on the non-profit's homepage, technical reports, and platform pages — *tokens donated, frameworks contributed, benchmark tasks authored, skills released, compute supplied, studies co-designed*. Where for-profit partners generate revenue through the platform, those revenues can be routed back as donations and redistributed to skill-contributing developers.

Who is building it

ACADEMIC BACKING

Stanford · Princeton · Oxford · UIUC · Berkeley

FOUNDING TEAM

Researchers who have led or co-led the largest agent-focused workshops in the field. Partnerships with leading open-source agent communities and top-tier agent startups.

WHY NOW

The defaults for AI-assisted science are being set now. Workflows, benchmarks, and interfaces standardized over the next one to two years will be difficult to displace. A credible open alternative has a short window in which to be established — and OpenScientists exists to use it.